

# Music Genre Classification

Davide Pietrasanta - 844824


Giuseppe Magazzù - 829612

Gaetano Magazzù - 829685

# Task

The task for this project is the **classification** of music genres from **audio** files.

It is assumed that only one musical genre is associated with each audio track.



# Dataset

## Free Music Archive (FMA)

- Raw audio tracks + additional metadata
- Suggested split (Train 80%, Validation 10%, Test 10%)
- Available subsets:
  - **small: 8000 audio tracks and 8 musical genres (balanced)**
  - medium: 25000 audio tracks and 16 musical genres (unbalanced)
  - large: 106,574 audio tracks and 161 musical genres (unbalanced)
- Pre-computed features: statistical moments for different spectral features

# Workflow

The work followed the following development:

- Handcrafted features
- CNN
  - Simple CNN
  - Hyper-parameters optimization
  - Data augmentation
  - Transfer Learning



# Handcrafted features

## FFNN


Input: 518 pre-computed features

Output: 8 genres classes

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	265728
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 8)	2056

=====  
Total params: 399,112  
Trainable params: 399,112  
Non-trainable params: 0



Early Stopping - Best Epoch: 35

### Inference Time

CPU: 20.70ms  
GPU: 20.07ms

Train	loss: 1.470, accuracy: 0.680
Validation	loss: 1.715, accuracy: 0.567
Test	loss: 1.900, accuracy: <b>0.495</b>

# Feature for CNN

## Pre-processing:

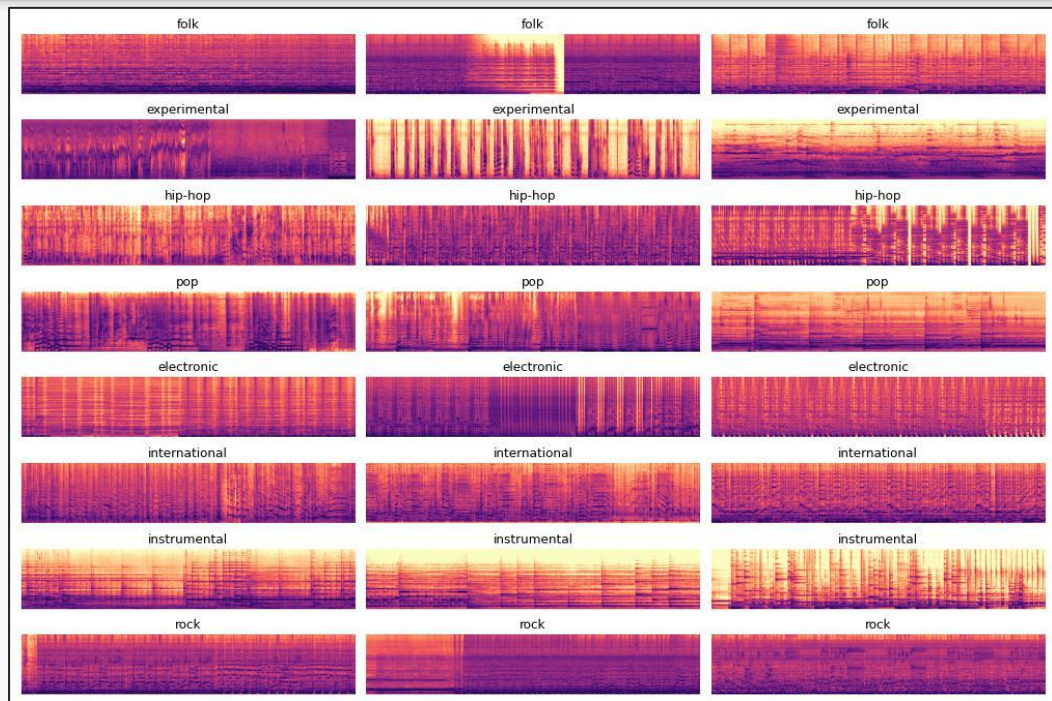
- Resample at 22050 Hz
- Trim audio at 29.70 seconds
- Removed corrupted audio files (6)

## Log Mel Spectrogram as Images:

- Grayscale [0, 255]
- (128, 1280)

## Intraclass Variation

dataset: <https://www.kaggle.com/giuseppemagazz/fma-mel-new>



# Simple CNN

The architecture used was the one suggested by "Daniel Kostrzewa".

The number of trainable parameters is 735,944.


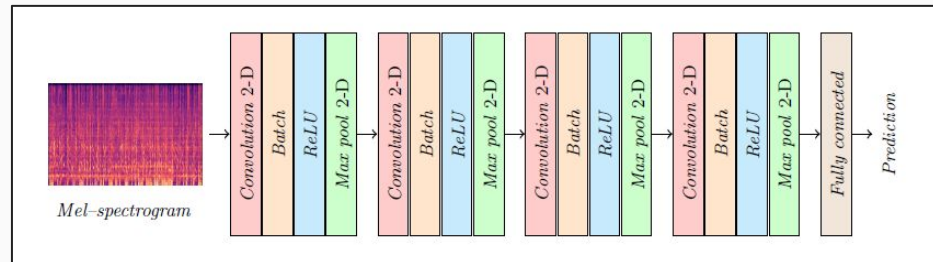
The model reaches an accuracy of **0.3975** on the test set and, as can be seen from the graphs, it overfits.

The accuracy of the paper is, however, equal to **0.5163**.

Time (CPU): 64 ms

Time(GPU): 42 ms

D. Kostrzewa, P. Kaminski, and R. Brzeski, "Music genre classification: Looking for the perfect network," in International Conference on Computational Science. Springer, 2021, pp. 55–67.



# Hyper-parameters optimization

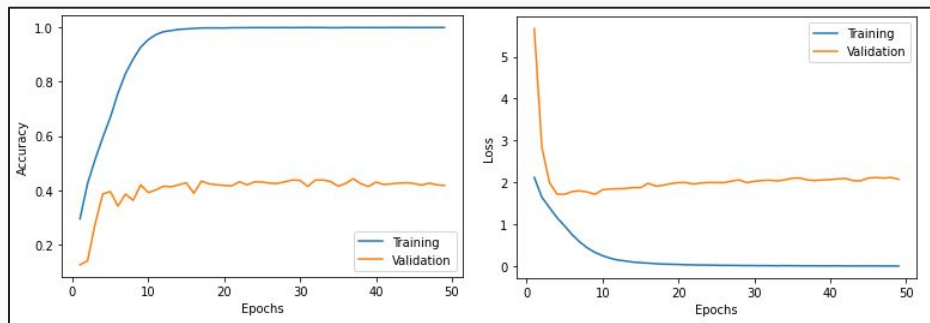
We then perform a hyper-parameter optimization using Hyperband.

The hyper-parameters considered are:

- The size of the kernel (fixed on each layer) in [3,6].
- The number of kernels for each layer in [32, 256].
- The probability of dropout in {0.2, 0.25, 0.5}.
- The initial learning rate in {0.01, 0.001, 0.0001}.

The model reaches an accuracy of **0.3900** on the test set and, as can be seen from the graphs, it overfits.

We can see how the validation curve is more stable this time.



Time (CPU): 60 ms

Time (GPU): 43 ms



# Data augmentation

As data augmentation we have used SpecAugment:

- Frequency Masking
- Time Masking


Two new examples for each example in the train set (Train: 19196, Validation: 800, Test: 800).

After the data augmentation the new accuracies for the test set are:

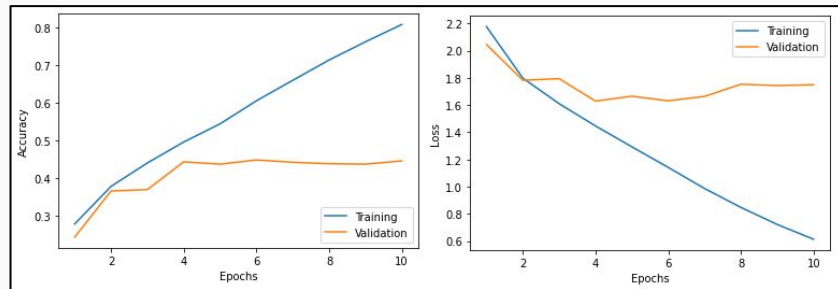
Simple CNN: **0.3225**

Tuned CNN: **0.3787**

We still have overfitting and in the case of the simple CNN we have really unstable results on the validation.



Above, the curves of the simple CNN. Below, the curves of the Tuned CNN




# Transfer Learning

We performed features extraction on two Top-5 Accuracy CNN on ImageNet VGG16 e ResNet50.

Low amount of data for training

Tasks really different



# Experiments

New split training/test with a test of 20%, (6398, 1596).

3 classifiers were trained by a 10-fold cross validation with a grid search for hyperparameters tuning.

Different percentages of cumulative explained variance were tested on FC2 layer of VGG16: 90%, 95%, 99%.

Cut Level	Size	Size after PCA	PCA value	Model	Accuracy		
					svm linear	svm rbf	mlp
Fc2	4096	1480	0.99	VGG16	0.4440/0.4003	0.5168/0.4454	0.4795/0.4191
Fc2	4096	429	0.95	VGG16	0.4965/0.4530	0.5206/0.4542	0.4890/0.4361
Fc2	4096	153	0.9	VGG16	0.5099/0.4586	0.5195/0.4617	0.4920/0.4398

Train/Test Accuracy

# Experiments

## VGG16


Cut Level	Size	Size after PCA	Model	Accuracy		
				svm linear	svm rbf	mlp
Fc2	4096	153	VGG16	0.5099/0.4586	0.5195/0.4617	0.4920/0.4398
Fc1	4096	311	VGG16	0.5157/0.4573	0.5315/0.4837	0.5076/0.4718
block5_pool	7x7x512	1174	VGG16	0.4316/0.4185	<b>0.5317/ 0.4968</b>	0.4726/0.4429

## ResNet50

Cut Level	Size	Size after PCA	Model	Accuracy		
				svm linear	svm rbf	mlp
avg_pool	1000	87	ResNet50	0.5268/0.4649	0.5381/0.4874	0.5246/0.4592
conv5_block1_2_relu	7x7x512	1321	ResNet50	0.4193/0.4104	<b>0.5484/0.5012</b>	0.5221/0.4755

Train/Test Accuracy

# Comparison




VGG16 best cut + SVM RBF

**Accuracy:** 0.4968

**FE:** 73.21ms(CPU)  
**FE:** 23.90ms (GPU)

**SVM:** 02.69ms (CPU)

**Parameters:** 14,714,688



FE = feature extraction

**Accuracy:** 0.5012

**FE:** 51.34ms (CPU)  
**FE:** 23.75ms (GPU)

**SVM:** 03.27ms (CPU)


**Parameters:** 11,477,888

ResNet50 best cut + SVM RBF

# Results

## Accuracies

- Handcrafted features: **0.495**
- Simple CNN, from paper: **0.5163**
- Simple CNN, our implementation: **0.3975**
- Transfer Learning, with ResNet50: **0.5012**



# Future developments

Possible ways to improve could be:

- Optimization of hyper-parameters from start to finish
- Consider medium and large dataset
- Different pre-processing
- Better data augmentation
- Other features as MFCCs and its derivatives



# Extra



# Handcrafted Features

For each of the following features are calculated:

- min
- max
- mean
- median
- std
- skew
- kurtosi

Total: 518 values

Features:

- Chroma STFT: 12
- Chroma CQT: 12
- Chroma CENS: 12
- Tonnezt: 6
- MFCC: 20
- Zero Crossing Rate: 1
- RMSE: 1
- Spectral Centroid: 1
- Spectral Bandwidth: 1
- Spectral Contrast: 7
- Spectral Rolloff: 1

# FFNN

Adam

Learning Rate: 1e-4

Batch: 32

Epochs: 200

Early Stopping (patience: 3)

Output Layer:

- Categorical Cross Entropy
- Softmax (8 classes)

Input Layer:

- 518 features

Units	Parameters	Train Loss	Val Loss	Test Loss
(512, 512)	532,488	1.482	1.723	1.898
(512, 256)	399,112	1.470	1.715	1.900
(256, 256)	200,712	1.503	1.733	1.946
(256, 64)	149,832	1.475	1.733	1.943
(128, 64)	75,208	1.476	1.729	1.905
(64)	33,736	1.332	1.598	1.790
(128)	67,464	1.321	1.577	1.800
(256)	134,920	1.338	1.592	1.806
(512)	269,832	1.334	1.579	1.803



Table 1: Values of loss for the different number of neurons.

Units	Parameters	Train Acc	Val Acc	Test Acc
(512, 512)	532,488	0.677	0.563	0.487
(512, 256)	399,112	0.680	0.567	<b>0.495</b>
(256, 256)	200,712	0.683	0.577	0.482
(256, 64)	149,832	0.690	0.577	0.484
(128, 64)	75,208	0.683	0.564	0.489
(64)	33,736	0.680	0.560	0.465
(128)	67,464	0.686	0.566	0.475
(256)	134,920	0.685	0.566	0.479
(512)	269,832	0.683	0.571	0.472


Table 2: Values of accuracy for the different number of neurons.

# Data Augmentation (SpegAugment)

Frequency Masking



Time Masking



The pixel range is randomly generated from a uniform distribution.  
For the frequency (0, 27) was used, while for the time (0, 100).


# Extra: Hyperband

**Hyperband** extends the **Successive Halving** algorithm.

The idea behind the original Successive Halving algorithm is about uniformly allocate a budget to a set of hyper-parameter configurations, evaluate the performance of all configurations, throw out the worst half, and repeat until one configuration remains.

It requires the **number of configurations**  $n$  and some finite **budget**  $B$  as inputs to the algorithm.

Then  $B/n$  resources are allocated on average across the configurations.



# Extra: Hyperband

However, in many cases it's not known if it's better to consider many configurations or not.

It runs several Successive Halving runs with different budgets and number of configurations, to find the best set.

It begins with the maximum exploration to ends up with a classical random search, in which every configuration is allocated with  $R$  resources.

**Algorithm 1:** HYPERBAND algorithm for hyperparameter optimization.


```
input      :  $R, \eta$  (default  $\eta = 3$ )
initialization:  $s_{\max} = \lfloor \log_{\eta}(R) \rfloor, B = (s_{\max} + 1)R$ 
1 for  $s \in \{s_{\max}, s_{\max} - 1, \dots, 0\}$  do
2    $n = \lceil \frac{B}{R} \frac{\eta^s}{(s+1)} \rceil, \quad r = R\eta^{-s}$ 
   // begin SUCCESSIVEHALVING with  $(n, r)$  inner loop
3    $T = \text{get\_hyperparameter\_configuration}(n)$ 
4   for  $i \in \{0, \dots, s\}$  do
5      $n_i = \lfloor n\eta^{-i} \rfloor$ 
6      $r_i = r\eta^i$ 
7      $L = \{\text{run\_then\_return\_val\_loss}(t, r_i) : t \in T\}$ 
8      $T = \text{top.k}(T, L, \lfloor n_i/\eta \rfloor)$ 
9   end
10 end
```

11 return Configuration with the smallest intermediate loss seen so far.

L. Li, K. Jamieson, G. De Salvo, R. A. Talwalkar, and A. Hyperband, "A novel bandit-based approach to hyperparameter optimization," Computer Vision and Pattern Recognition, arXiv: 1603.0656, 2016

# Extra: Parameters after tuning

	Simple CNN	Tuned CNN	Tuned CNN (data augmentation)
Layer 1 (num kernels)	64	32	64
Layer 2 (num kernels)	64	128	64
Layer 3 (num kernels)	128	192	64
Layer 4 (num kernels)	128	98	96
kernel size	5	6	4
Dropout	0,2	0,25	0,25
Learning rate	0,001	0,001	0,0001



# EXTRA: ML classifiers tuning

**linear svm:** C: {0.01, 0.1, 1, 10}

**svm rbf:** C: {0.01, 0.1, 1, 10}, gamma: {0.01, 0.001, 0.000}

**mlp:**

- optimizer=adam
- learning\_rate=1e-4
- batch\_size=32
- max\_epochs=200 with early stopping
- activations all relu
- last layer softmax(8 classes)
- cross entropy loss

**network architectures:** [(512, 256), (512, 32), (512,)]

**l2 regularization alpha:** [0.01, 0.03, 0.05]

N	svm linear	svm rbf	mlp
FC2, PCA=0.99	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.03, 'hidden_layer_sizes': (512, 32)}
FC2, PCA=0.95	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.05, 'hidden_layer_sizes': (512, 32)}
FC2, PCA=0.90	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.05, 'hidden_layer_sizes': (512, 256)}

Table 8: hyperparameters result for the VGG16 FC2 with different values of pca.

Cut Level	svm linear	svm rbf	mlp
avg_pool	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.05, 'hidden_layer_sizes': (512, 32)}
conv5_block1_2_relu	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.01, 'hidden_layer_sizes': (512, )}

Table 12: hyperparameters result for different cut of ResNet50 with PCA=0.90.

Cut Level	svm linear	svm rbf	mlp
Fc2	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.05, 'hidden_layer_sizes': (512, 256)}
Fc1	{'C': 0.01}	{'C': 10, 'gamma': 0.0001}	{'alpha': 0.05, 'hidden_layer_sizes': (512, 32)}
block5_pool	{'C': 0.01}	{'C': 1, 'gamma': 0.0001}	{'alpha': 0.03, 'hidden_layer_sizes': (512, 256)}

Table 10: hyperparameters result for different cut level of VGG16 with PCA=0.90.